

MitoChime: A Machine Learning Pipeline for Detecting PCR-Induced Chimeras in Mitochondrial Illumina Reads

Yvonne Lin¹, Duranne Duran¹, Daniella Pailden¹, and Francis Dimzon¹

University of the Philippines Visayas, Miagao, Iloilo, Philippines
linyvonnea@gmail.com, duranduranne@gmail.com, daniellalpailden@gmail.com,
fddimzon1@up.edu.ph

Abstract. Polymerase chain reaction (PCR) amplification during Illumina library preparation can introduce chimeric reads that degrade mitochondrial genome assembly by creating spurious junctions and fragmented assembly graphs. Existing chimera detection tools are largely optimized for amplicon-based microbial workflows and often rely on reference databases or abundance assumptions that do not fit single-species organellar sequencing. We present a pipeline for detecting PCR-induced chimeric reads in simulated *Sardinella lemuru* mitochondrial paired-end data using alignment-derived features and a sequence-based convolutional model. Using simulated clean and chimeric reads, alignment structure and clipping metrics are extracted alongside k-mer discontinuity and microhomology descriptors as features. On a held-out test set, a Gradient Boosting classifier achieves F1=0.777 and ROC-AUC=0.846 for chimeric-read detection, while a 1D Convolutional Neural Network model achieves F1=0.914 and ROC-AUC=0.961 with higher recall. Assembly evaluation using SPAdes-generated assembly graphs summarized with Bandage statistics shows that filtering reduced assembly graph complexity at high chimera rates, with performance depending on read depth. These results support MitoChime as a practical pre-assembly quality-control step for mitochondrial genome assembly pipelines.

Keywords: Chimera detection · Mitochondrial genome assembly · Illumina reads · Machine learning · Deep learning

1 Introduction

Polymerase chain reaction (PCR) amplification during next-generation sequencing (NGS) library preparation can generate chimeric reads, which are artificial recombinant sequences formed when partially copied DNA fragments attach to different templates and are extended in later cycles [10,13]. As shown in Fig. 1, fragments from two templates can combine and be amplified into one artificial read. Despite the high throughput and accuracy of Illumina sequencing, PCR-induced chimeras remain a persistent artifact that can affect downstream analyses and compromise genome assembly. Reported PCR-induced chimera rates

vary across studies and experimental conditions, particularly in amplicon-based datasets, indicating that chimera formation is protocol-dependent rather than fixed [13].




Fig. 1: Schematic illustration of PCR-induced chimera formation through incomplete extension and template switching, adapted from [16].

Mitochondrial genomes are particularly susceptible to this issue because they are small, circular, and often contain repetitive regions [2,4]. Under some conditions, junction-supporting artifacts may introduce misleading connections in assembly graphs, produce branching structures or fragmentation, or in some cases contribute to assembly failure [5,8,9]. Because many organelle assembly tools such as SPAdes [1], GetOrganelle [9], and MITObim [8] assume that input reads are largely free of such artifacts, chimera-related assembly failure modes remain insufficiently addressed.

Many existing chimera detection approaches were originally developed for amplicon-based microbial community sequencing and often use differences in sequence frequency to help identify chimeras [7,15]. Although effective in many microbial studies, these approaches are not well suited to single-species organellar sequencing, highlighting the need for a more specialized method. A pipeline that detects chimeric reads from alignment and sequence irregularities can therefore serve as a useful pre-assembly filtering step.

Sardinella lemuru was selected as the target organism because it is an economically important sardine species in the Philippines and a relevant subject for mitochondrial genome analysis [11,17]. Its selection was also supported by the availability of mitochondrial assemblies and raw sequencing datasets in public repositories such as the National Center for Biotechnology Information (NCBI), which facilitates reference selection and benchmarking. This study introduces MitoChime, a pipeline for detecting PCR-induced chimeric reads in simulated *Sardinella lemuru* mitochondrial Illumina paired-end data. The pipeline extracts interpretable features from standard BAM and FASTQ artifacts, benchmarks su-

pervised machine learning models together with a sequence-based convolutional neural network (CNN), and evaluates the effects of filtering on downstream assembly graph complexity.

2 Related Work

Chimera detection in NGS has been widely studied in amplicon sequencing, where tools such as UCHIME and UCHIME2 detect sequences that appear to be recombined from two parent sequences [7,6]. DADA2 detects chimeras in amplicon data at the level of inferred sequence variants rather than from irregularities in individual reads [3]. Although effective in many microbial studies, these approaches are not designed for read-level detection of PCR-induced chimeras in single-species organellar sequencing data.

Ensemble approaches such as CATCh use the scores and binary decisions produced by multiple chimera detection tools as input to a supervised learning model, improving detection in 16S rRNA amplicon sequencing but still relying on the sequence and alignment signals examined by those tools [12]. Pipelines such as ChimPipe detect biological chimeras in RNA-seq using discordant paired-end and split-read evidence, but are intended for fusion detection and usually require reference genome and annotation data [14].

In contrast, mitochondrial assembly pipelines generally do not provide explicit read-level chimera filtering tailored to PCR artifacts. This gap motivates a feature-driven approach that targets alignment fragmentation and sequence discontinuity in mitochondrial paired-end reads and evaluates success in terms of downstream assembly graph outcomes.

3 Methods

3.1 Overview

Figure 2 summarizes the MitoChime workflow: paired-end read simulation, alignment to a mitochondrial reference, (i) tabular feature extraction for classical machine learning and (ii) sequence-only dataset preparation for deep learning, followed by model inference, pair-safe read filtering, and downstream assembly evaluation.

3.2 Data generation and alignment

Paired-end Illumina reads were simulated from the *Sardinella lemuru* mitochondrial reference (NC_039553.1). A total of 40,000 paired-end reads were generated, comprising 20,000 chimeric and 20,000 clean read pairs. Clean read pairs were generated with **wgsim** (v1.20) at 150 bp using an error rate of 0.05. Chimeric templates were constructed by concatenating two non-adjacent reference segments, optionally retaining short overlaps to approximate PCR template switching, and were simulated with the same read-length configuration but with




Fig. 2: MitoChime workflow from simulated reads to filtering and assembly evaluation.

the error rate set to 0 to preserve junction signals. Reads were aligned to the reference using `minimap2` (v2.30-r1287) in short-read mode (`-ax sr`) and processed with `samtools` (v1.20) to produce sorted and indexed BAM files. An illustrative example of one read from the dataset is shown in Figure 3.

Example data point from the dataset

```

read_id: NC_039553.1_10001_10492_5:0:0_8:0:0_b77/1
label: 0
seq: GTAAGGCATAGTAAGGCAATGGTGAGGCTGTGTGTATTGCAGCAGTTCAAATTCAC...

```

Fig. 3: Illustrative example of one record from the simulated dataset. The sequence is truncated for readability.

3.3 Feature and sequence extraction

Two complementary input representations were prepared from the same labeled read pairs to support both classical machine learning and deep learning.

The extracted features included: (i) alignment-structure summaries, such as the presence and number of supplementary alignments, distances between aligned segments, strand consistency, mapping quality, and mismatch counts; (ii) clipping features derived from the CIGAR string, including left and right soft clipping, total clipped bases, and an estimated breakpoint position; (iii) sequence-composition discontinuity between the left and right portions of the read, measured using 6-mer profiles; and (iv) microhomology features around

the estimated breakpoint, including the longest short overlap and its GC content.

In parallel, for deep learning, the raw nucleotide sequence (length $L = 150$) was retained as a separate input for the CNN. Sequences were one-hot encoded at training and inference time, with four base channels (A, C, G, T). Ambiguous bases were represented as all zeros.

3.4 Modeling and evaluation protocol

The labeled dataset was partitioned into train and test splits (80/20) using stratification to preserve the balanced class ratio. Model selection, tuning, and five-fold cross validation were performed on the training split only. Final performance was reported on the held-out test split using accuracy, precision, recall, F1, and ROC-AUC (precision/recall/F1 for the chimeric class).

Classical machine learning (feature-based). Feature-based models were implemented in Python using `scikit-learn`. An initial benchmark evaluated a panel of classifiers under five-fold stratified cross-validation on the training split, including a dummy baseline, L2-regularized logistic regression, a calibrated linear SVM, k-nearest neighbours, Gaussian Naïve Bayes, tree ensembles (Random Forest, ExtraTrees, and bagging with decision trees), and gradient boosting methods (Gradient Boosting (GB), XGBoost, LightGBM, and CatBoost). Based on cross-validation ranking, a subset of competitive model families was carried forward for hyperparameter optimization. Hyperparameters were tuned using randomized search with five-fold stratified cross-validation (`RandomizedSearchCV`), optimizing chimeric-class F1. Among the feature-based models, a tuned Gradient Boosting classifier achieved the best validation ranking and was used as the primary classical model. The hyperparameter search included the number of estimators, and `learning_rate` (selected value 0.2); the selected model was refit on the full training split and evaluated once on the held-out test split.

Deep learning (sequence-based). A one-dimensional convolutional neural network (CNN1D) was trained using sequence-only input of length $L = 150$, represented by one-hot encoding. The architecture comprised three convolutional blocks with 64, 128, and 256 filters (kernel sizes 7, 5, and 3; ReLU activations; max-pooling after the first two blocks), followed by adaptive max pooling and a two-layer MLP head ($256 \rightarrow 128 \rightarrow 2$) with dropout (0.2). Training used mini-batches of 128 for up to 30 epochs with learning rate 0.001 and weight decay 10^{-4} . The selected checkpoint was the model achieving the best held-out F1 during training.

3.5 Filtering strategy

Predicted chimeric reads were removed using a probability threshold τ on the model’s chimeric-class score. For each read pair, a chimeric probability p was assigned, and the pair was retained if $p < \tau$. Pair consistency was enforced by

filtering on a normalized base read identifier, ensuring that both mates were retained or removed together in the output FASTQ files.

A threshold of $\tau = 0.5$ was used as the baseline operating point for assembly evaluation because it represents the standard decision boundary for binary classifiers and allows the Gradient Boosting and CNN1D models to be compared under the same filtering rule. This threshold was not optimized separately for each read depth or chimera rate. Therefore, the results should be interpreted as a baseline comparison of model behavior rather than as an optimized filtering strategy.

3.6 Assembly and downstream evaluation

Assemblies were generated from unfiltered and filtered read sets using **SPAdes** (v4.2.0). Assembly graphs were summarized using **Bandage** (v0.8.1) graph statistics. Graph complexity was quantified using graph-level metrics including node and total assembled length, across simulated read depths and chimera rates.

3.7 Chimeric Detection Pipeline

A command-line interface automates preprocessing, model inference, pair-safe filtering, and optional assembly evaluation (see Fig. 4) to support reproducible use in mitochondrial assembly workflows.




Fig. 4: MitoChime pipeline for model selection, filtering, and optional assembly evaluation.

4 Results

4.1 Held-out test performance

Table 1 summarizes held-out test performance for the tuned GB (feature-based) classifier and the CNN1D (sequence-based) model. The tuned GB model exhibits high precision (0.937) but lower recall (0.663), indicating a conservative decision strategy that reduced false positives at the expense of missed chimeric reads. In contrast, the CNN1D substantially increases recall (0.959) and overall discrimination (ROC-AUC 0.961), at the cost of reduced precision (0.872), resulting in more clean reads flagged as chimeric.

Table 1: Held-out test performance of the tuned Gradient Boosting model and CNN1D. Precision/recall/F1 are reported for the chimeric (positive) class.

Model	Accuracy	Precision	Recall	F1	ROC-AUC
Gradient Boosting (tuned)	0.809	0.937	0.663	0.777	0.846
CNN1D (final)	0.909	0.872	0.959	0.912	0.961

4.2 Error profiles

The trade-off between the two models is shown in Fig. 5. GB produces relatively few false positives (179 clean reads predicted as chimeric) but many false negatives (1,348 chimeric reads predicted as clean). CNN1D reverses this pattern, with markedly fewer false negatives (165) but more false positives (561). These complementary error profiles motivate evaluating both models in terms of downstream assembly outcomes rather than classification metrics alone.




Fig. 5: Confusion matrices for the Gradient Boosting and CNN1D models on the held-out test.

4.3 External Validation Under Varying Read Depths and Chimera Rates

Table 2 reports **Bandage** and **SPAdes** graph statistics for unfiltered and filtered assemblies across read depths and chimera rates. In this evaluation, fewer graph nodes generally indicate a simpler and more contiguous mitochondrial assembly, whereas many nodes suggest fragmentation or branching caused by conflicting read evidence. Filtering substantially reduces graph complexity under high chimera rates. At 20k read pairs, unfiltered assemblies were already simple at 5% chimeras (1 node), and this structure was preserved under CNN1D filtering, whereas GB produced a slightly more complex graph (2 nodes). At 10–15% chimeras, unfiltered graphs increased to 3 nodes, while both GB and CNN1D filtering collapsed the graph to 1 node. Under the highest chimeric percentage (50% chimeras) shown in Fig. 6, the unfiltered assembly became highly fragmented (966 nodes, total length 163,906 bp). Both filters substantially reduced fragmentation, with GB producing 3 nodes and CNN1D restoring a 1-node graph.

At 10k read pairs, assemblies remained 1 node for 5–15% chimeras under all conditions, indicating limited graph-level disruption at moderate chimera rates. At 50% chimeras, unfiltered graphs increased to 3 nodes, while both GB and CNN reduced complexity to 2 nodes.

At 3.2k read pairs, both filtering methods simplified the 5% chimera case from 2 nodes (unfiltered) to 1 node. However, under 15% chimeras, CNN filtering increased fragmentation relative to both unfiltered and GB (3 nodes), and at 50% chimeras, shown in Fig. 7, CNN produced 2 nodes, whereas unfiltered and GB remained at 1 node.




Fig. 6: **Bandage** assembly graphs for 20k reads with 50% chimeras at a 50% filtering threshold.

Table 2: Assembly statistics for unfiltered and filtered reads across read depths and chimera rates. Nodes summarize graph complexity. Reads kept indicates the percentage of input reads retained after filtering. Max length is reported from **SPAdes** assemblies.

Depth	Chimera%	Condition	Nodes	Reads kept (%)	Max length
3.2k	5	Unfiltered	2	100.000	16599
	5	GB	1	97.406	16577
	5	CNN	1	82.125	16575
	10	Unfiltered	1	100.000	16591
	10	GB	1	95.219	16615
	10	CNN	1	79.812	16590
	15	Unfiltered	1	100.000	16645
	15	GB	1	93.375	16606
	15	CNN	3	75.188	14925
	50	Unfiltered	1	100.000	16736
	50	GB	1	77.875	16646
	50	CNN	2	48.781	16247
10k	5	Unfiltered	1	100.000	16615
	5	GB	1	97.320	16615
	5	CNN	1	83.090	16603
	10	Unfiltered	1	100.000	16704
	10	GB	1	95.300	16614
	10	CNN	1	78.470	16693
	15	Unfiltered	1	100.000	16735
	15	GB	1	93.330	16642
	15	CNN	1	75.950	16611
	50	Unfiltered	3	100.000	16715
	50	GB	2	78.940	16787
	50	CNN	2	49.470	16709
20k	5	Unfiltered	1	100.000	16669
	5	GB	2	97.490	16619
	5	CNN	1	82.575	16669
	10	Unfiltered	3	100.000	16799
	10	GB	1	95.135	16616
	10	CNN	1	78.935	16616
	15	Unfiltered	3	100.000	16747
	15	GB	1	93.345	16891
	15	CNN	1	75.220	16646
	50	Unfiltered	966	100.000	407
	50	GB	3	78.485	16799
	50	CNN	1	49.890	16745



Fig. 7: Bandage assembly graphs for 3200 reads with 50% chimeras at a 50% filtering threshold.

5 Discussion

Results show that PCR-induced chimeras in short-read mitochondrial data can be detected and effectively removed using interpretable features derived from standard alignments, and that a sequence-based CNN1D can further improve sensitivity.

The observed assembly behavior is consistent with the classification profiles of the two models. Gradient Boosting achieved higher precision but lower recall, indicating a more conservative filtering strategy that removed fewer clean reads at the cost of leaving more chimeric reads behind. This explains why Gradient Boosting performed more robustly at low read depth, where preserving true reads is critical for maintaining assembly support. In contrast, CNN1D achieved much higher recall but lower precision, indicating a more aggressive filtering strategy that removed more chimeric reads but also misclassified more clean reads as chimeric. This trade-off appears advantageous at high read depth, particularly in the 20K-read, 50% chimera condition, where CNN1D fully restored a 1-node graph while Gradient Boosting left residual branching. However, at 3200 reads, the lower precision of CNN1D resulted in over-filtering, removing enough true reads to fragment the graph despite improved chimera detection. Because the present evaluation used simulated reads from a single mitochondrial reference, performance on empirical datasets may differ. Real sequencing libraries can contain platform-specific errors, uneven coverage, nuclear mitochondrial segments, contaminants, and library-preparation artifacts that were not fully represented in the simulation. Different taxa may also vary in mitochondrial genome structure, repeat content, GC composition, and divergence from the selected reference, all of which may affect alignment patterns and sequence-discontinuity features. Likewise, library preparation protocols that use different polymerases, cycle numbers, or insert-size distributions may produce different chimera profiles. These factors may change the precision–recall balance of both models and may require threshold calibration before deployment on new datasets.

6 Conclusion

MitoChime offers a practical pre-assembly chimera detection pipeline for simulated mitochondrial Illumina paired-end reads using interpretable alignment-derived and sequence-based signals. The tuned Gradient Boosting classifier achieves robust performance with conservative filtering, while a CNN1D model provides higher sensitivity and stronger overall chimera detection. External assembly evaluation further shows that filtering can substantially reduce assembly graph complexity under high chimera percentage, with coverage-dependent trade-offs at low depth. In general, the Gradient Boosting model is better suited to low-coverage datasets where preserving read support is critical, while the CNN1D is more advantageous at moderate-to-high coverage where its aggressive chimera detection can be fully utilized without risking assembly fragmentation. These trade-offs may also be adjusted in the chimera detection pipeline through model selection

and threshold tuning, depending on the characteristics of the input data and the specific goals of the user. Overall, these findings support MitoChime as a useful quality-control step for improving mitochondrial genome assembly workflows, particularly in resource-constrained settings.

The main limitations of this study are its reliance on simulated reads and evaluation on a single target species. Although simulation enabled controlled benchmarking with known clean and chimeric labels, it cannot fully capture the complexity of empirical sequencing libraries, including uneven coverage, real sequencing errors, contamination, and protocol-specific artifacts. Broader validation across taxa, library preparation conditions, and real mitochondrial datasets is therefore required to establish generalizability. Future work will focus on calibration using empirical sequencing libraries, evaluation under varied sequencing error profiles, and support for additional references and model families.

Code Availability The source code, configuration files, and scripts used for simulated read generation, chimera construction, alignment, feature extraction, model training, filtering, and assembly evaluation are available at:

<https://github.com/linyvonnea/MitoChimeV1>.

Acknowledgements The authors thank the Bioinformatics team of the Philippine Genome Center Visayas for technical assistance and helpful discussions during this study. The authors also acknowledge the support of DOST-SEI and DOST Region VI through the scholarship grants awarded to the authors.

References

1. Bankevich, A., Nurk, S., Antipov, D., Gurevich, A.A., Dvorkin, M., Kulikov, A.S., Lesin, V.M., Nikolenko, S.I., Pham, S., Prjibelski, A.D., Pyshkin, A.V., Sirotkin, A.V., Vyahhi, N., Tesler, G., Alekseyev, M.A., Pevzner, P.A.: SPAdes: A new genome assembly algorithm and its applications to single-cell sequencing. *Journal of Computational Biology* **19**(5), 455–477 (2012). <https://doi.org/10.1089/cmb.2012.0021>
2. Boore, J.L.: Animal mitochondrial genomes. *Nucleic Acids Research* **27**(8), 1767–1780 (1999). <https://doi.org/10.1093/nar/27.8.1767>
3. Callahan, B.J., McMurdie, P.J., Rosen, M.J., Han, A.W., Johnson, A.J.A., Holmes, S.P.: DADA2: High-resolution sample inference from illumina amplicon data. *Nature Methods* **13**(7), 581–583 (2016). <https://doi.org/10.1038/nmeth.3869>
4. Cameron, S.L.: Insect mitochondrial genomics: Implications for evolution and phylogeny. *Annual Review of Entomology* **59**, 95–117 (2014). <https://doi.org/10.1146/annurev-ento-011613-162007>
5. Dierckxsens, N., Mardulyn, P., Smits, G.: NOVOPlasty: de novo assembly of organelle genomes from whole genome data. *Nucleic Acids Research* **45**(4), e18 (2017). <https://doi.org/10.1093/nar/gkw955>
6. Edgar, R.C.: UCHIME2: improved chimera prediction for amplicon sequencing. *BioRxiv* (2016), <https://api.semanticscholar.org/CorpusID:88955007>

7. Edgar, R.C., Haas, B.J., Clemente, J.C., Quince, C., Knight, R.: UCHIME improves sensitivity and speed of chimera detection. *Bioinformatics* **27**(16), 2194–2200 (2011). <https://doi.org/10.1093/bioinformatics/btr381>
8. Hahn, C., Bachmann, L., Chevreux, B.: Reconstructing mitochondrial genomes directly from genomic next-generation sequencing reads—a baiting and iterative mapping approach. *Nucleic Acids Research* **41**(13), e129 (2013). <https://doi.org/10.1093/nar/gkt371>
9. Jin, J.J., Yu, W.B., Yang, J., Song, Y., dePamphilis, C.W., Yi, T.S., Li, D.Z.: GetOrganelle: a fast and versatile toolkit for accurate de novo assembly of organelle genomes. *Genome Biology* **21**(1), 241 (2020). <https://doi.org/10.1186/s13059-020-02154-5>
10. Judo, M.S.B., Wedel, W.R., Wilson, B.H.: Stimulation and suppression of PCR-mediated recombination. *Nucleic Acids Research* **26**(7), 1819–1825 (1998). <https://doi.org/10.1093/nar/26.7.1819>
11. Labrador, K., Agmata, A., Palermo, J.D., Ravago-Gotanco, R., Pante, M.J.: Mitochondrial dna reveals genetically structured haplogroups of bali sardinella (*Sardinella lemuru*) in philippine waters. *Regional Studies in Marine Science* **41**, 101588 (2021). <https://doi.org/10.1016/j.rsma.2020.101588>
12. Mysara, M., Saeys, Y., Leys, N., Raes, J., Monsieurs, P.: CATCh, an ensemble classifier for chimera detection in 16s rRNA sequencing studies. *Applied and Environmental Microbiology* **81**(5), 1573–1584 (2015). <https://doi.org/10.1128/AEM.02896-14>, <https://journals.asm.org/doi/abs/10.1128/AEM.02896-14>
13. Qiu, X., Wu, L., Huang, H., McDonel, P.E., Palumbo, A.V., Tiedje, J.M., Zhou, J.: Evaluation of PCR-generated chimeras, mutations, and heteroduplexes with 16s rRNA gene-based cloning. *Applied and Environmental Microbiology* **67**(2), 880–887 (2001). <https://doi.org/10.1128/AEM.67.2.880-887.2001>
14. Rodriguez-Martin, B., Palumbo, E., Marco-Sola, S., Griebel, T., Ribeca, P., Alonso, G., Djebali, S.: ChimPipe: Accurate detection of fusion genes and transcription-induced chimeras from RNA-seq data. *BMC Genomics* **18** (2017). <https://doi.org/10.1186/s12864-016-3404-9>
15. Rognes, T., Flouri, T., Nichols, B., Quince, C., Mahé, F.: VSEARCH: a versatile open source tool for metagenomics. *PeerJ* **4**, e2584 (2016). <https://doi.org/10.7717/peerj.2584>
16. Sieber, F.: Development of a Tool to Address Nucleic Acids into Mitochondria. Ph.D. thesis, Nantes Université (Sep 2011), PhD thesis, advisor: Laurence Maréchal-Drouard
17. Willette, D., Bognot, E., Mutia, M.T., Santos, M.: Biology and ecology of sardines in the philippines: A review. Tech. Rep. 1, NFRDI Technical Paper Series (2011)